

# Team 39

# Predict Closed Questions on StackOverflow

...

November 1, 2016

Rounak Mundra, Sourav Chanduka, Arihant Jain

# Problem Statement


- "Millions of programmers use StackOverflow to get high quality answers to their programming questions every day.
- More than six thousand new questions is asked on StackOverflow every weekday. Currently about 6% of all new questions end up "closed".

**Goal:-** It is to build a classifier that predicts whether or not a question will be closed given the question is submitted.

# Dataset Information

- **Number of Instances** : 0.4 Million (135 MB in CSV Format)
- **Data Available in Dataset** :
  - PostId
  - PostCreationDate
  - OwnerCreationDate
  - ReputationAtPostCreation
  - OwnerUndeletedAnswerCountAtPostTime
  - Title
  - BodyMarkdown
  - 5 Tags
  - PostClosedDate
  - **Status (Open/Closed) (0/1)**
- **Source of Dataset** : Kaggle.com

# Feature Selection





# Features Extraction

- Status: (0/1) (Open/Closed)
- TitleLength (After text processing)
- BodyLength
- ReputationAtPostCreation: Reputation of User when he posted a question
- NumberOfTags (atleast 1 and upto 5)
- CloseCreationDateDifference: Number of days it took to get deleted.
- UserAge: Profile Age
- OwnerUndeletedAnswerCountAtPostTime
- LinksCount


# Text Processing (Preprocessing)


- Both Title and Body contains text.
- We need to re-present the text in numbers without losing the relevance of it.
- Natural Language Process Techniques:
  - **Tokenizing**: converting a document to its atomic elements.
  - **Stopping**: removing meaningless words.
  - **Stemming**: merging words that are equivalent in meaning.
- After these processing we got 10% more accurate prediction than when we took length straight away.

















# Algorithms Implemented

## Random Forest

Random forests are a combination of tree predictors such that each tree depends on the values of a random vector sampled independently and with the same distribution for all trees in the forest.

## Support Vector Machine

Support Vector Machines are based on the concept of decision planes that define decision boundaries. A decision plane is one that separates between a set of objects having different class memberships.

## Vowpal Wabbit

The Vowpal Wabbit (VW) project is a fast out-of-core learning system sponsored by Microsoft Research and (previously) Yahoo! Research.

# Random Forest

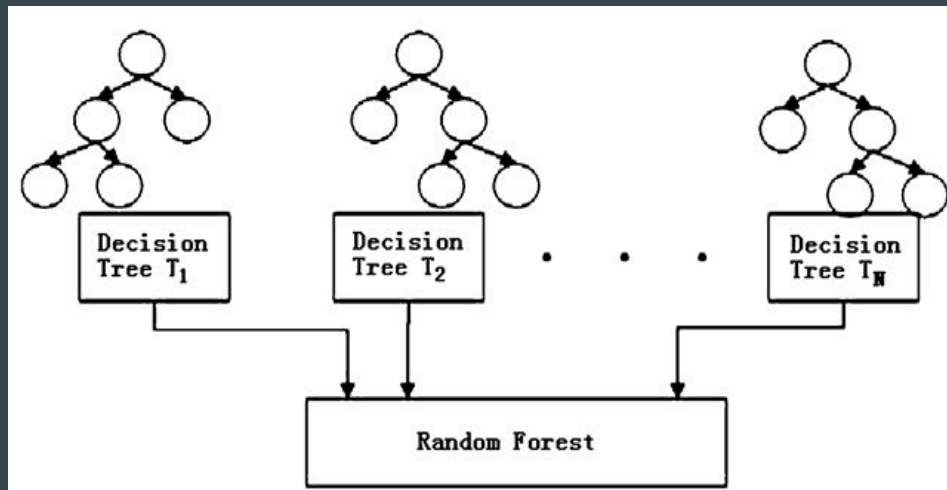
No. of Instances predicted: 520

Accuracy : 65 - 75 %

Precision : 0.65 - 0.70

F1 Score : 0.55 - 0.65

Recall : 0.75 - 0.85



# Support Vector Machine (SVM)


No. of Instances predicted: 520

Accuracy : Around 65 %

Precision : 0.60 - 0.65

F1 Score : 0.50 - 0.60

Recall : 0.77 - 0.82



# Vowpal Wabbit

1. Vowpal Wabbit is a modified stochastic gradient descend algorithm.
2. Pre-processing CSV files to get CSV files slightly more fitting our purpose.
3. VW allows to induce sparsity in learned feature weights
4. For Text processing, it has it's own online Latent Dirichlet Allocation (LDA) implementation.
5. VW focuses on the approach to stream the examples to an online learning algorithm in contrast of parallelization of a batch learning algorithm over many machines.
6. **Average Loss** is the value of loss function as the learner goes along with example counter.
7. The main idea of truncate gradient is that it uses the simple rounding rule of weight to achieve the sparsity



# Training Analysis

```
sourav@tesla:~/SMAI_Project$ python extract.py train-tiny.csv train-tiny.csv
sourav@tesla:~/SMAI_Project$ python csv2vw.py train-tiny.csv train.vw
sourav@tesla:~/SMAI_Project$ vw --loss_function logistic --oaa 5 -d train.vw -f model
final_regressor = model
Num weight bits = 18
learning rate = 0.5
initial_t = 0
power_t = 0.5
using no cache
Reading datafile = train.vw
num sources = 1
```

average loss	since last	example counter	example weight	current label	current predict	current features
0.333333	0.333333	3	3.0	4	4	34
0.333333	0.333333	6	6.0	4	4	71
0.454545	0.600000	11	11.0	4	4	37
0.454545	0.454545	22	22.0	3	4	41
0.431818	0.409091	44	44.0	2	4	50
0.436782	0.441860	87	87.0	4	4	61
0.442529	0.448276	174	174.0	1	4	30
0.456897	0.471264	348	348.0	1	4	29

```
finished run
number of examples = 537
weighted example sum = 537
weighted label sum = 0
average loss = 0.476723
best constant = 0
total feature number = 43552
sourav@tesla:~/SMAI_Project$
```

# Results And Prediction

Efficiency :  $1 - \text{Average loss} = 1 - 0.137 = 0.863$

```
sourav@tesla:~/kaggle-stackoverflow$ vw --loss_function logistic --oaa 5 -i model -t -d train-tiny.vw -r raw_predictions.txt
only testing
Num weight bits = 18
learning rate = 10
initial_t = 1
power_t = 0.5
raw predictions = raw_predictions.txt
using no cache
Reading datafile = train-tiny.vw
num sources = 1
```

average loss	since last	example counter	example weight	current label	current predict	current features
0.000000	0.000000	3	3.0	4	4	34
0.000000	0.000000	6	6.0	4	4	71
0.090909	0.200000	11	11.0	4	4	37
0.090909	0.090909	22	22.0	3	3	41
0.068182	0.045455	44	44.0	2	2	50
0.045977	0.023256	87	87.0	4	4	61
0.063218	0.080460	174	174.0	1	1	30
0.091954	0.120690	348	348.0	1	1	29

```
finished run
number of examples = 537
weighted example sum = 537
weighted label sum = 0
average loss = 0.137803
best constant = -0.00186567
total feature number = 43552
sourav@tesla:~/kaggle-stackoverflow$
```

**Thank You !**