Team 39
Predict Closed Questions on StackOverflow

• • •

November 1, 2016
Rounak Mundra, Sourav Chanduka, Arihant Jain

Problem Statement

- "Millions of programmers use StackOverflow to get high quality answers to their programming questions every day.
- More than six thousand new questions is asked on StackOverflow every weekday. Currently about 6% of all new questions end up "closed".

Goal:- It is to build a classifier that predicts whether or not a question will be closed given the question as submitted.

Dataset Information

- **Number of Instances** : 0.12 Million (135 MB in CSV Format)
- **Data Available in Dataset** :
    - PostId
    - PostCreationDate
    - OwnerCreationDate
    - ReputationAtPostCreation
    - OwnerUndeletedAnswerCountAtPostTime
    - Title
    - BodyMarkdown
    - 5 Tags
    - PostClosedDate
    - **Status (Open/Closed) (0/1)**
- **Source of Dataset :** Kaggle.com

Feature Selection

Features Selection

- Status: (0/1) (Open/Closed)
- TitleLength (After text processing)
- BodyLength
- ReputationAtPostCreation: Reputation of User when he posted a question
- NumberOfTags (atleast 1 and upto 5)
- CloseCreationDateDifference: Number of days it took to get deleted.
- UserAge: Profile Age
- OwnerUndeletedAnswerCountAtPostTime
- LinksCount

Text Processing

- Both Title and Body contains text.
- We need to re-present the text in numbers without losing the relevance of it.
- Natural Language Process Techniques:
  - **Tokenizing**: converting a document to its atomic elements.
  - **Stopping**: removing meaningless words.
  - **Stemming**: merging words that are equivalent in meaning.
- After these processing we got 10% more accurate prediction than when we took length straight away.

Algorithms Implemented

Random Forest

Random forests are a combination of tree predictors such that each tree depends on the values of a random vector sampled independently and with the same distribution for all trees in the forest.

Support Vector Machine

Support Vector Machines are based on the concept of decision planes that define decision boundaries. A decision plane is one that separates between a set of objects having different class memberships.

Vowpal Wabbit

The Vowpal Wabbit (VW) project is a fast out-of-core learning system sponsored by Microsoft Research and (previously) Yahoo! Research.

Random Forest

No. of Instances predicted: 520

Accuracy : 65 - 75 %

Precision : 0.65 - 0.70

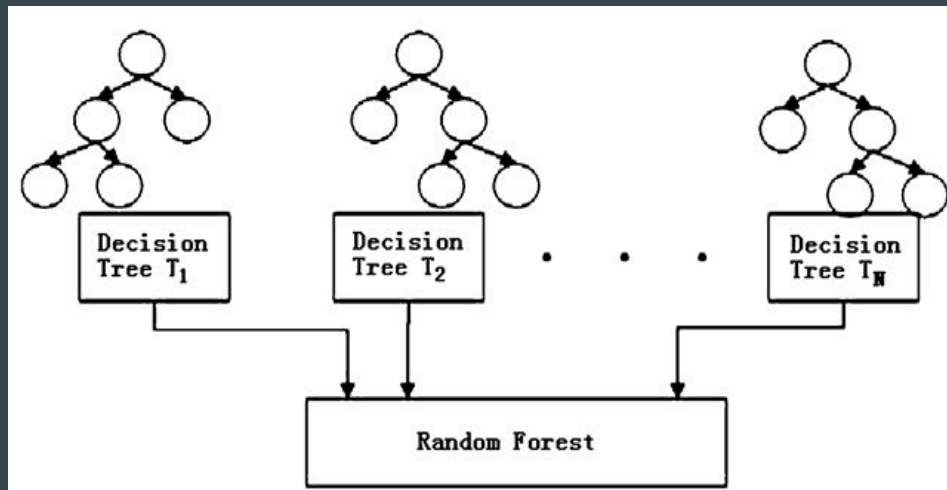F1 Score : 0.55 - 0.65

Recall : 0.75 - 0.85

Support Vector Machine (SVM)

No. of Instances predicted: 520

Accuracy : Around 65 %

Precision : 0.60 - 0.65

F1 Score : 0.50 - 0.60

Recall : 0.77 - 0.82

SVM Algorithm

Class 2

Support Vectors

Optimal Separating Hyperplane

Class 1

Finding the Optimal Separating Hyperplane in SVM

Vowpal Wabbit

1. Vowpal Wabbit is a modified stochastic gradient descend algorithm.
2. Pre-processing CSV files to get CSV files slightly more fitting our purpose.
3. VW allows to induce sparsity in learned feature weights
4. For Text processing, it has it's own online Latent Dirichlet Allocation (LDA) implementation.
5. VW focuses on the approach to stream the examples to an online learning algorithm in contrast of parallelization of a batch learning algorithm over many machines.
6. **Average Loss** is the value of loss function as the learner goes along with example counter.
7. The main idea of truncate gradient is that it uses the simple rounding rule of weight to achieve the sparsity

Training Analysis

```
sourav@tesla:~/SMAI_Project$ python extract.py train-tiny.csv train-tiny_.csv
sourav@tesla:~/SMAI_Project$ python csv2vw.py train-tiny_.csv train.vw
sourav@tesla:~/SMAI_Project$ vw --loss_function logistic --oaa 5 -d train.vw -f model
final_regressor = model
Num weight bits = 18
learning rate = 0.5
initial_t = 0
power_t = 0.5
using no cache
Reading datafile = train.vw
num sources = 1
average      since           example     example   current  current  current
loss         last            counter      weight     label  predict features
0.333333     0.333333              3         3.0         4        4       34
0.333333     0.333333              6         6.0         4        4       71
0.454545     0.600000             11        11.0         4        4       37
0.454545     0.454545             22        22.0         3        4       41
0.431818     0.409091             44        44.0         2        4       50
0.436782     0.441860             87        87.0         4        4       61
0.442529     0.448276            174       174.0         1        4       30
0.456897     0.471264            348       348.0         1        4       29

finished run
number of examples = 537
weighted example sum = 537
weighted label sum = 0
average loss = 0.476723
best constant = 0
total feature number = 43552
sourav@tesla:~/SMAI_Project$ ▮
```

Results And Prediction

Efficiency : 1 - Average loss = 1 - 0.137 = 0.863

```
sourav@tesla:~/kaggle-stackoverflow$ vw --loss_function logistic --oaa 5 -i model -t -d train-tiny.vw -r raw_predictions.txt
only testing
Num weight bits = 18
learning rate = 10
initial_t = 1
power_t = 0.5
raw predictions = raw_predictions.txt
using no cache
Reading datafile = train-tiny.vw
num sources = 1
average     since        example    example   current  current  current
loss        last         counter     weight    label    predict  features
0.000000    0.000000           3       3.0        4        4        34
0.000000    0.000000           6       6.0        4        4        71
0.090909    0.200000          11      11.0        4        4        37
0.090909    0.090909          22      22.0        3        3        41
0.068182    0.045455          44      44.0        2        2        50
0.045977    0.023256          87      87.0        4        4        61
0.063218    0.080460         174     174.0        1        1        30
0.091954    0.120690         348     348.0        1        1        29

finished run
number of examples = 537
weighted example sum = 537
weighted label sum = 0
average loss = 0.137803
best constant = -0.00186567
total feature number = 43552
sourav@tesla:~/kaggle-stackoverflow$
```

Thank You !